


Note: All code and plots can be found [here!](#)

Problem 1.

Solution:

- a) The surface plot and code to produce it are included below. There appear to be four local minima of the function on this domain.



Code to generate the plot is included below:

```
1 ##problem 1
2 function f(vec::Vector{Float64})
3     #break vector input into x and y components
4     x = vec[1]
5     y = vec[2]
6     #evaluate function
7     val =(x^2 + y - 11)^2 + (x + y^2 - 7)^2
8     val
9 end
10 #create domain grid and matrix for z values
11 x_grid = collect(-4.0:0.01: 4.0)
12 nx = length(x_grid)
13 z_grid = zeros(nx,nx)
```

```
14 #loop over x and y grids
15 for i=1:nx, j=1:nx
16     #evaluate function at corresponding grid values and set z_grid
    equal to function value
17     z_grid[i,j] = f(x_grid[i], x_grid[j])
18 end
19
20 #plot and save figure
21 plot_1a = Plots.surface(x_grid, x_grid, z_grid, seriescolor=:viridis,
    camera=(25,70))
22 savefig(plot_1a, "plot1a.png")
```

b) The gradient is given by

$$\nabla(x, y) = \langle 4x(x^2 + y - 11) + 2(x + y^2 - 7), 2(x^2 + y - 11) + 4y(x + y^2 - 7) \rangle$$

and the Hessian is given by

$$H(x, y) = \begin{bmatrix} 12x^2 + 4y - 42 & 4x + 4y \\ 4x + 4y & 4x + 12y^2 - 26 \end{bmatrix}$$

Using Newton's method, we find that the four local minima are located at (-3.78, -3.28), (3.58, -1.85), (3,2), and (-2.81, 3.13). These are obtained using initial guesses (-4,-4), (4,-4), (4,4), and (-4, 4). Each of these guesses is located near one of the local minima, so the algorithm ends up finding the closest one to each initial guess.

```
1 #create gradient function
2 function g(G, guess::Vector{Float64})
3     #split guess vector into component parts
4     x,y = guess[1], guess[2]
5     G[1] = 4*x*(x^2 + y - 11) + 2*(x + y^2 - 7)
6     G[2] = 2*(x^2+y-11) + 4*y*(x + y^2 - 7)
```

```
7      G
8  end
9  #create Hessian function
10 function h(H, guess::Vector{Float64})
11     #split guess vector into component parts
12     x,y = guess[1], guess[2]
13     #define hessian matrix values
14     H[1,1] = 12*x^2 + 4*y - 42
15     H[1,2] = 4x + 4y
16     H[2,1] = 4x + 4y
17     H[2,2] = 4x + 12y^2 - 26
18     H
19 end
20
21 using Optim
22 #create guesses (should have done a loop)
23 guess1 = [-4.0,-4.0]
24 guess2 = [4.0,-4.0]
25 guess3 = [4.0,4.0]
26 guess4 = [-4.0,4.0]
27
28 #find minima given initial guesses
29 opt1 = optimize(f, g,h, guess1)
30 opt2 = optimize(f, g,h, guess2)
31 opt3 = optimize(f, g,h, guess3)
32 opt4 = optimize(f, g,h, guess4)
33 println(opt1.minimizer)
34 println(opt2.minimizer)
35 println(opt3.minimizer)
36 println(opt4.minimizer)
37 >>[-3.7793102534670995, -3.2831859913022914]
38 >>[3.584428340395532, -1.84812652689675]
39 >>[3.000000000059233, 1.999999999921049]
```

```
40 >>[-2.805118087105429, 3.131312518242223]
```

c) For $(-3.77, -3.28)$, Nedler-Mead took 38 iterations and Newton's method took 96 iterations. For $(3.58, -1.85)$, Nedler-Mead took 34 iterations and Newton's method took 72. For $(3,2)$, Nedler-Mead took 39 iterations and Newton's method took 13 iterations. For $(-2.81, 3.13)$, Nedler-Mead took 42 iterations and Newton's method took 758 (!) iterations. Code is included below:

```
1 opt1n = optimize(f, guess1)
2 opt2n = optimize(f, guess2)
3 opt3n = optimize(f, guess3)
4 opt4n = optimize(f, guess4)
```


Nedler-Mead offers better performance for all of the minima except for $(3,2)$.


□

Problem 2.

Solution:

a) Surface and contour plots are included below. Upon inspection of the graphs, we find that the global minimum is located at $(0,0)$.





```
1 #Ackley function
2 function Ackley(vec::Vector{Float64})
3     #take vector and split into x and y components
4     x,y = vec[1], vec[2]
5     #evaluate function
6     val = -20*exp(-0.2*sqrt(0.5*(x^2+y^2))) - exp(0.5(cos(2*pi*x) +
7     cos(2*pi*y)))+ exp(1) + 20
8     val
9
10 #create x_grid and 2d grid for function values
11 x_grid = collect(-4.0:0.01: 4.0)
12 nx = length(x_grid)
13 z_grid = zeros(nx,nx)
14 #loop through x and y arrays and evaluate function at each point then
15     store in z_grid
16 for i=1:nx, j=1:nx
17     vec = [x_grid[i], x_grid[j]]
18     z_grid[i,j] = Ackley(vec)
19 end
20 #plot and save the surface and contour plots
```

```
21 plot_2a_surf = Plots.surface(x_grid, x_grid, z_grid, seriescolor=:  
    viridis, camera=(25,70))  
22 plot_2a_cont = Plots.contourf(x_grid, x_grid, z_grid, seriescolor=:  
    viridis)  
23  
24 savefig(plot_2a_surf, "plot2a_surf.png")  
25 savefig(plot_2a_cont, "plot2a_cont.png")  
26  
27
```

b) A table of initial guesses and algorithm performance is included below. Upon inspection, it appears that LBFGS takes far fewer iterations than Nedler-Mead. However, both algorithms seem to miss the global minimum for values farther away from the origin. I chose some guesses that are close to the local minima not at the origin, and it appears that both algorithms can get tripped up at different places. Code to generate the guesses is included below as well.

Guess	NM Min	NM Iter	LBFGS Min	LBFGS Iter
(0.1, 0.1)	(0, 0)	51	(0, 0)	9
(2.0, 1.0)	(1.96, 0.98)	32	(-0.95, 0)	8
(2.0, 1.5)	(0, 0)	57	(0.98, 1.96)	9
(0, 1.0)	(0, 0.95)	25	(0, 0)	9
(2.0, 2.0)	(0, 0)	63	(0.97, 0.97)	4
(2.2, 2.2)	(0, 0)	58	(0, 0)	11
(3.0, 3.2)	(2.96, -1.17)	40	(1.17, 0)	20

Table 1: Guesses, minimizers, and iterations

```
1 #initialize array of guesses to try  
2 guesses = [[0.1, 0.1], [2,1], [2, 1.5], [0, 1], [2,2], [2.2,2.2],  
    [3,3.2]]  
3 #loop through guesses in array  
4 for guess in guesses  
5     println(guess)  
6     #run both optimization algorithms
```

```
7  opt_nm = optimize(Ackley, guess)
8  opt_lb = optimize(Ackley, guess, LBFGS())
9  #print results
10 println("NM",opt_nm.minimizer, opt_nm.iterations)
11 println("LBFGS",opt_lb.minimizer, opt_lb.iterations)
12 end
```

□


Problem 3.

Solution:

- a) The global minimum is zero at $x = 0$. The minimum holds for arbitrary n . Note that $x_i^2 \geq 0$ and $-A \cos(2\pi x_i) \geq -A \forall x_i \in \mathbb{R}$. Hence, we have that

$$f(x) = An + \sum_{i=1}^n [x_i^2 - A \cos(2\pi x_i)] \geq An - \sum_{i=1}^n A = An - An = 0 = f(0)$$

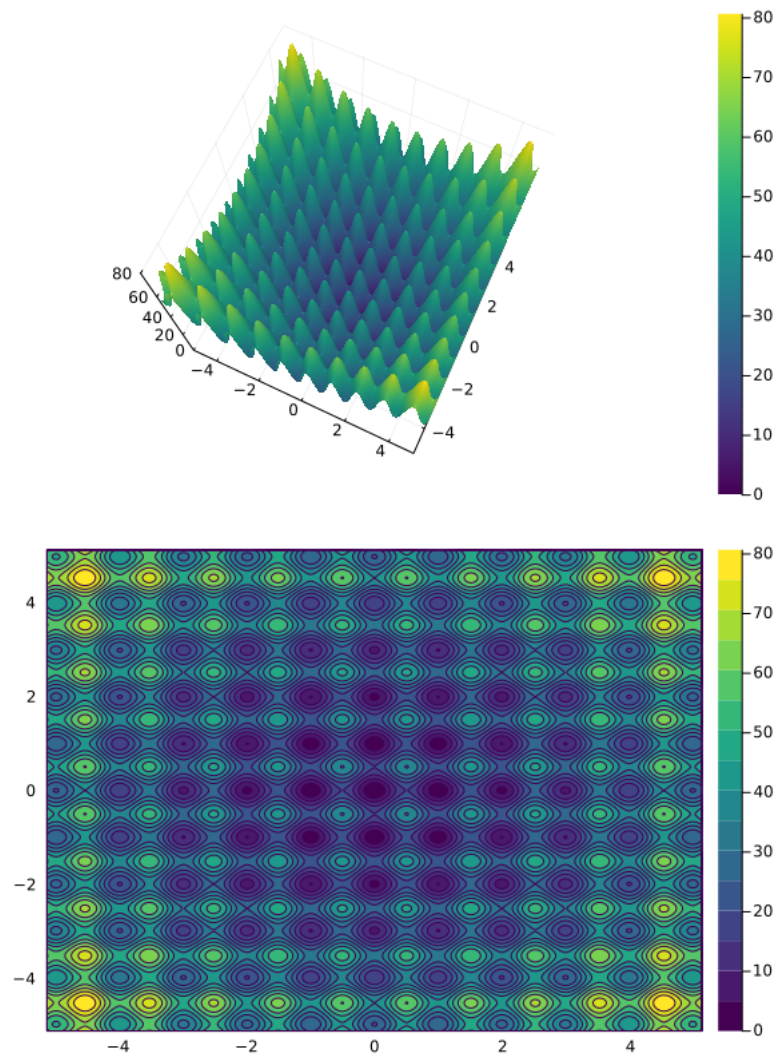
where we can see that the function attains its global minimum at the origin for any arbitrary n . A plot of the function and the code used to generate it are included below:



```
1 function Rastrigin(vec::Vector{Float64})
```

```
2     #start with the constant term
3     val = 10*length(vec)
4     #summation term: iteratively add the x_i terms
5     for x_i in vec
6         val += x_i^2 - 10*cos(2*pi*x_i)
7     end
8     val
9 end
10
11 #create arrays for domain and function values
12 x_grid = collect(-5.12:0.01:5.12)
13 nx = length(x_grid)
14 z_grid = zeros(nx)
15 #loop over domain and evaluate function at each point
16 for i=1:nx
17     z_grid[i] = Rastrigin([x_grid[i]])
18 end
19 using Plots
20 #plot and save figure
21 plot_3a = plot(x_grid, z_grid)
22 savefig(plot_3a, "plot3a.png")
```

b) The plots and code for $n = 2$ are included below:



```
1 #create array of zeros for plotting
2 z_grid = zeros(nx,nx)
3 #loop over x and y and evaluate function
4 for i=1:nx, j = 1:nx
5     #create vector to feed into the function
6     vec = [x_grid[i], x_grid[j]]
7     #evaluate function
8     z_grid[i,j] = Rastrigin(vec)
9 end
10
11 #plot and save the surface and contour plots
```

```
12 plot_3b_surf = Plots.surface(x_grid, x_grid, z_grid, seriescolor=:  
    viridis, camera=(25,70))  
13 plot_3b_cont = Plots.contourf(x_grid, x_grid, z_grid, seriescolor=:  
    viridis)  
14 savefig(plot_3b_surf, "plot3asurf.png")  
15 savefig(plot_3b_cont, "plot3acont.png")
```

c) A table of initial guesses and algorithm performance is included below. For these guesses, LBFGS finds the global minimum far more consistently than does Nedler-Mead. Furthermore, LBFGS is significantly faster. Code used to generate the table is included below.

Guess	NM Min	NM Iter	LBFGS Min	LBFGS Iter
(0.1, 0.1)	(0, 0)	32	(0, 0)	3
(0.5, 0.5)	(0.99, 0.99)	35	(0, 0)	2
(1.0, 1.0)	(0.99, 0.99)	34	(0, 0)	2
(1.0, 0)	(1, 0)	29	(0, 0)	2
(0, 1.0)	(0, 1.0)	29	(0, 0)	2
(0, 0.5)	(0, 1)	32	(0, 0)	2
(0.5, 0)	(1.0, 0)	32	(0, 0)	2
(2.0, 2.0)	(0, 0)	43	(0, 0)	2
(2.2, 2.2)	(0, 0)	43	(-0.99, -0.99)	4
(3.0, 3.2)	(2.98, 3.98)	35	(3.0, 3.2)	1

Table 2: Guesses, minimizers, and iterations

```
1 #initialize array of guesses to try  
2 guesses = [[0.1, 0.1], [0.5,0.5], [1,1],[1, 0], [0,1], [0,0.5],  
    [0.5,0], [2.0, 2.0], [2.2,2.2], [3.0, 3.2]]  
3  
4 #loop through guesses in array  
5 for guess in guesses  
6     println(guess)  
7     #run both optimization algorithms  
8     opt_nm = optimize(Rastrigin, guess)  
9     opt_lb = optimize(Rastrigin, guess, LBFGS())
```

```
10     #print results
11     println("NM",opt_nm.minimizer, opt_nm.iterations)
12     println("LBFGS",opt_lb.minimizer, opt_lb.iterations, "\n")
13 end
```

□

Problem 4.

Solution: Code (and output) for problem 4 is included below. I tested my interpolation function for $f(x) = x^2$ on the interval $[0, 2]$ with 10 grid points at the point $x = 1.3$. The interpolated value is 1.696, whereas the true value of the function at that point is 1.69.

```
1 function next_highest(x_arr::Vector{Float64}, x::Float64)
2     #create boolean array where 1 indicates x is less than y and 0 if y is
3     #less than x
4     compare_arr = [isless(x,y) for y in x_arr]
5     #find index of first y such that x < y
6     i_first = findfirst(compare_arr)
7     #find the corresponding value in x_arr
8     x_first = x_arr[i_first]
9     x_first, Int(i_first)
10 end
11
12 function lin_approx(f, a::Float64, b::Float64, n::Int64, x::Float64)
13     #check if x is lower than a, return a
14     if x <= a
15         val = f(a)
16         return val
17     #likewise for b
18     elseif x >= b
19         val = f(b)
20         return val
21     else #for values of x in the interval
```


```
21     #create domain and empty array for the function values
22     x_grid = collect(range(a, length = n , stop= b))
23     #fill in the function values
24     z_grid = f.(x_grid)
25     #find the lowest value in the x grid which is bigger than x, as
    well as its index
26     x_high, i_high = next_highest(x_grid, x)
27     #since x_high is the smallest value which is greater than x, the
    element which comes before it must be less than or equal to x
28     x_low = x_grid[Int(i_high-1)]
29     #find the weighted average of x_high and x_low to get the
    interpolated value
30     x_interp = ((x-x_low)/(x_high - x_low))*f(x_high) + ((x_high-x)/(x
    _high - x_low))*f(x_low)
31     return x_interp
32 end
33 end
34
35 f(x) = x^2
36
37 #find interpolation for f at x=1.3 on interval [0,2] with 10 grid points
38 lin_approx(f, 0.0, 2.0, 10, 1.3)
39 >> 1.696
```

□

Problem 5.

Solution:

- a) Using an evenly-spaced grid, I generated the following plot of the linear interpolation, the actual function values, and the interpolation error. Code is included below the figure:



```
1 f(x) = log(x+1)
2 function next_highest(x_arr::Vector{Float64}, x::Float64)
3     #create boolean array where 1 indicates x is less than y and 0 if
4     #y is less than x
5     compare_arr = [isless(x,y) for y in x_arr]
6     #find index of first y such that x < y
7     i_first = findfirst(compare_arr)
8     #find the corresponding value in x_arr
9     x_first = x_arr[i_first]
10    x_first, Int(i_first)
11 end
12 function lin_approx(f, x_ar::Vector{Float64}, x::Float64)
13     #sort array so that it is ordered
14     x_grid = sort(x_ar)
15     #check if x lies outside of the boundaries of the grid
16     if x <= x_grid[1] #if x is lower than the lowest point in the
17         grid:
18         val = f(x_grid[1])
19         return val
20     elseif x >= x_grid[length(x_grid)] #if x is higher than the
21         highest point in the grid:
```

```
20     val = f(x_grid[length(x_grid)])
21     return val
22 else #if x is within the grid
23     z_grid = f.(x_grid) #fill in the function values at each
point of the x_grid
24     x_high, i_high = next_highest(x_grid, x) #find the value in
the grid just above x
25     x_low = x_grid[Int(i_high-1)] #find the value in the grid
just below x
26     #find the interpolated value using x_low and x_high
27     x_interp = ((x-x_low)/(x_high - x_low))*f(x_high) + ((x_high-
x)/(x_high - x_low))*f(x_low)
28     return x_interp
29 end
30 end
31
32 function interp_eval(f, x_arr::Vector{Float64})
33     #create fine x grid and find its length
34     x_fine = collect(0.0:0.01:100.0)
35     nx = length(x_fine)
36     #create empty grid for interpolated values on the fine grid
37     z_fine_approx = zeros(nx)
38     #find the actual function values for each point on the fine grid
39     z_fine_act = f.(x_fine)
40     #loop over x_grid
41     for i=1:nx
42         #find the interpolated values for each point and record them
in z_fine_approx
43         z_fine_approx[i] = lin_approx(f, x_arr, x_fine[i])
44     end
45     #find the pointwise interpolation error
46     error = z_fine_act - z_fine_approx
47     #return the interpolated values, true values, interpolation error
```

```
    , and the fine grid used
48     z_fine_approx, z_fine_act, error, x_fine
49 end
50
51 #create evenly spaced grid
52 x_even = [0.0, 10.0, 20.0, 30.0, 40.0, 50.0, 60.0, 70.0, 80.0, 90.0,
           100.0]
53 #find interpolation of f using the above grid and store the required
    variables for plotting
54 z_fine_approx, z_fine_act, err, x_fine = interp_eval(f, x_even)
55
56 using Plots
57 #create and save plot
58 plot_5a = plot(x_fine, [z_fine_approx, z_fine_act, err], labels=["
    Interpolation" "Actual" "Error"], title="Linear interpolation, f(x
    ), and error", xlabel="x", ylabel="y")
59 savefig(plot_5a, "plot5a.png")
60
```

- b) Below is my function to take an arbitrary array and return the interpolation error. I chose an arbitrary array to test it and reported the sum of errors for that grid.

```
1 function interp_eval_arb(x_arr::Vector{Float64})
2     #add 0 and 100 and sort the resulting array so that it is ordered
3     x_arr_s = sort( union( 0.0, x_arr, 100.0))
4     #for purposes of optimization in part c, return a very high error
    value if any values of input array are out of bounds
5     if x_arr_s[1] < 0.0 || x_arr_s[11] > 100.0
6         return 100000
7     else
8         #create x array
9         x_fine = collect(0.0:0.01:100.0)
10        nx = length(x_fine)
```

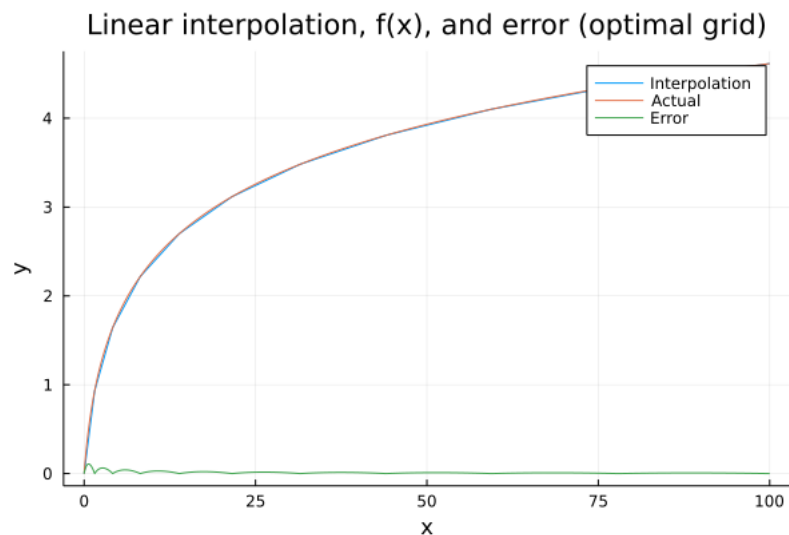
```
11     #initialize blank array for interpolated values
12     z_fine_approx = zeros(nx)
13     #evaluate function for grid points in fine grid
14     z_fine_act = f.(x_fine)
15     #loop over fine array
16     for i=1:nx
17         #find interpolated values for each point in array
18         z_fine_approx[i] = lin_approx(f, x_arr_s, x_fine[i])
19     end
20     #find interpolation error
21     error = z_fine_act - z_fine_approx
22     #find and return sum of the absolute value of interpolation
    error
23     sum_error = sum(abs.(error))
24     return sum_error
25 end
26 end
27
28 x_arb = [5.0, 10.0, 15.0, 20.0, 25.0, 30.0, 35.0, 40.0, 90.0]
29 #find interpolation error using the above arbitrary array
30 se = interp_eval_arb(x_arb)
31 >>419.07
32
```

- c) Below is the result of minimizing the interpolation error using Nedler-Mead. The optimal grid is (0, 1.54, 4.17, 8.19, 13.89, 21.57, 31.52, 44.05, 59.44, 77.99, 100.0). The final error is 110.355, which is down from 504.491 for the evenly-spaced grid. I also include a similar plot to that in problem 5a), and it is clear that this selection of grid points allows the interpolation to fit the actual function far better.

```
1 using Optim
2 x_guess = [10.0, 20.0, 30.0, 40.0, 50.0, 60.0, 70.0, 80.0, 90.0]
3 #find interpolation error using the above evenly-spaced starting
```



```
guess
4 se_guess = interp_eval_arb(x_guess)
5 >>504.491
6 #use Nedler-Mead to find the optimal grid points for interpolation.
  Our function rejects out-of-bounds guesses by giving a large
  penalty, and it also sorts the input array so we don't have to
  worry about checking that the points are in the right order
7 opt = optimize(interp_eval_arb, x_guess)
8 #find the optimal grid points from the optimization object
9 opt_grid = sort(opt.minimizer)
10 #find the sum of errors associated with optimal grid points
11 se_opt = interp_eval_arb(opt_grid)
12 >>110.355
13
```



Code for plot:

```
1 #function modified to return the required data for plotting
2 function interp_eval_arb_plot(x_arr::Vector{Float64})
3     x_arr_s = sort( union( 0.0, x_arr, 100.0) )
4     if x_arr_s[1] < 0.0 || x_arr_s[11] > 100.0
```

```
5         return 100000
6     else
7         z_arr = f.(x_arr_s)
8         x_fine = collect(0.0:0.01:100.0)
9         nx = length(x_fine)
10        z_fine_approx = zeros(nx)
11        z_fine_act = f.(x_fine)
12        for i=1:nx
13            z_fine_approx[i] = lin_approx(f, x_arr_s, x_fine[i])
14        end
15        error = z_fine_act - z_fine_approx
16        sum_error = sum(abs.(error))
17        return z_fine_approx, z_fine_act, error, x_fine
18    end
19 end
20
21 z_fine_approx, z_fine_act, err, x_fine = interp_eval_arb_plot(opt_grid)
22 plot_5c = plot(x_fine, [z_fine_approx, z_fine_act, err], labels=["
    Interpolation" "Actual" "Error"], title="Linear interpolation, f(x),
    and error (optimal grid)", xlabel="x", ylabel="y")
23 savefig(plot_5c, "plot5c.png")
```

